

OpenCL Buffers and Complete Examples

A Collaboration Between
David Kaeli, Northeastern University
Benedict R. Gaster, AMD
© 2011

Instructor Notes

- This is a brief lecture which goes into some more details on OpenCL memory objects
 - Describes various flags that can be used to change how data is handled between host and device, like page-locked I/O and so on
- The aim of this lecture is to cover required OpenCL host code for buffer management and provide simple examples
- Code for context and buffer management discussed in examples in this lecture serves as templates for more complicated kernels
 - This allows the next 3 lectures to be focused solely on kernel optimizations like blocking, thread grouping and so on
- Examples covered
 - Simple image rotation example
 - Simple non-blocking matrix-matrix multiplication

Topics

- Using OpenCL buffers
 - Declaring buffers
 - Enqueue reading and writing of buffers
- Simple but complete examples
 - Image Rotation
 - Non-blocking Matrix Multiplication

Creating OpenCL Buffers

- Data used by OpenCL devices is stored in a “buffer” on the device
- An OpenCL buffer object is created using the following function

```
cl_mem bufferobj = clCreateBuffer (  
    cl_context context,           //Context name  
    cl_mem_flags flags,          //Memory flags  
    size_t size,                 //Memory size allocated in buffer  
    void *host_ptr,              //Host data  
    cl_int *errcode)             //Returned error code
```

- Data can implicitly be copied to the device using a host pointer parameter
 - In this case copy to device is invoked when kernel is enqueued

Memory Flags

- Memory flag field in `clCreateBuffer()` allows us to define characteristics of the buffer object

Memory Flag	Behavior
CL_MEM_READ_WRITE	Specifies memory read / write behavior
CL_MEM_WRITE_ONLY	
CL_MEM_READ_ONLY	
CL_MEM_USE_HOST_PTR	Implementations can cache the contents pointed to by <code>host_ptr</code> in device memory. This cached copy can be used when kernels are executed on a device.
CL_MEM_ALLOC_HOST_PTR	Specifies to the implementation to allocate memory from host accessible memory.
CL_MEM_COPY_HOST_PTR	Specifies to allocate memory for the object and copy the data from memory referenced by <code>host_ptr</code> .

Copying Buffers to Device

- `clEnqueueWriteBuffer()` is used to write a buffer object to device memory (from the host)
- Provides more control over copy process than using host pointer functionality of `clCreateBuffer()`
 - Allows waiting for events and blocking

```
cl_int clEnqueueWriteBuffer (  
    cl_command_queue queue,           //Command queue to device  
    cl_mem buffer,                   //OpenCL Buffer Object  
    cl_bool blocking_read,           //Blocking/Non-Blocking Flag  
    size_t offset,                   //Offset into buffer to write to  
    size_t cb,                       //Size of data  
    void *ptr,                       //Host pointer  
    cl_uint num_in_wait_list,        //Number of events in wait list  
    const cl_event * event_wait_list, //Array of events to wait for  
    cl_event *event)                 //Event handler for this function
```

Copying Buffers to Host

- `clEnqueueReadBuffer()` is used to read from a buffer object from device to host memory
- Similar to `clEnqueueWriteBuffer()`

```
cl_int clEnqueueReadBuffer (  
    cl_command_queue queue,           //Command queue to device  
    cl_mem buffer,                   //OpenCL Buffer Object  
    cl_bool blocking_read,           //Blocking/Non-Blocking Flag  
    size_t offset,                   //Offset to copy from  
    size_t cb,                       //Size of data  
    void *ptr,                       //Host pointer  
    cl_uint num_in_wait_list,        //Number of events in wait list  
    const cl_event * event_wait_list, //Array of events to wait for  
    cl_event *event)                 //Event handler for this function
```

- The vector addition example discussed in Lecture 2 and 3 provide simple code snippets for moving data to and from devices

Example 1 - Image Rotation

- A common image processing routine
 - Applications in matching, alignment, etc.
- New coordinates of point (x_1, y_1) when rotated by an angle Θ around (x_0, y_0)

$$x_2 = \cos(\theta) * (x_1 - x_0) - \sin(\theta) * (y_1 - y_0) + x_0$$

$$y_2 = \sin(\theta) * (x_1 - x_0) + \cos(\theta) * (y_1 - y_0) + y_0$$

- By rotating the image about the origin $(0,0)$ we get

$$x_2 = \cos(\theta) * (x_1) - \sin(\theta) * (y_1)$$

$$y_2 = \sin(\theta) * (x_1) + \cos(\theta) * (y_1)$$

- Each coordinate for every point in the image can be calculated independently

Original Image




Rotated Image (90°)



Image Rotation

- Input: To copy to device
 - Image (2D Matrix of floats)
 - Rotation parameters
 - Image dimensions
- Output: From device
 - Rotated Image
- Main Steps
 - Copy image to device by enqueueing a write to a buffer on the device from the host
 - Run the Image rotation kernel on input image
 - Copy output image to host by enqueueing a read from a buffer on the device



The OpenCL Kernel

- Parallel portion of the algorithm off-loaded to device
 - Most thought provoking part of coding process
- Steps to be done in Image Rotation kernel
 - Obtain coordinates of work item in work group
 - Read rotation parameters
 - Calculate destination coordinates
 - Read input and write rotated output at calculated coordinates
- Parallel kernel is not always this obvious.
 - Profiling of an application is often necessary to find the bottlenecks and locate the data parallelism
- In this example grid of output image decomposed into work items
 - Not all parts of the input image copied to the output image after rotation, corners of I/P image could be lost after rotation

OpenCL Kernel

```
__kernel void image_rotate(
    __global float * src_data, __global float * dest_data, //Data in global memory
    int W, int H, //Image Dimensions
    float sinTheta, float cosTheta) //Rotation Parameters
{
    //Thread gets its index within index space
    const int ix = get_global_id(0);
    const int iy = get_global_id(1);

    //Calculate location of data to move into ix and iy– Output decomposition as mentioned
    float xpos = ( ((float) ix)*cosTheta + ((float)iy)*sinTheta);
    float ypos = ( ((float) iy)*cosTheta - ((float)ix)*sinTheta);

    if ( ( ((int)xpos>=0) && ((int)xpos< W)) //Bound Checking
        && (((int)ypos>=0) && ((int)ypos< H)))
    {
        //Read (xpos,ypos) src_data and store at (ix,iy) in dest_data
        dest_data[iy*W+ix]=
            src_data[(int)(floor(ypos*W+xpos))];
    }
}
```

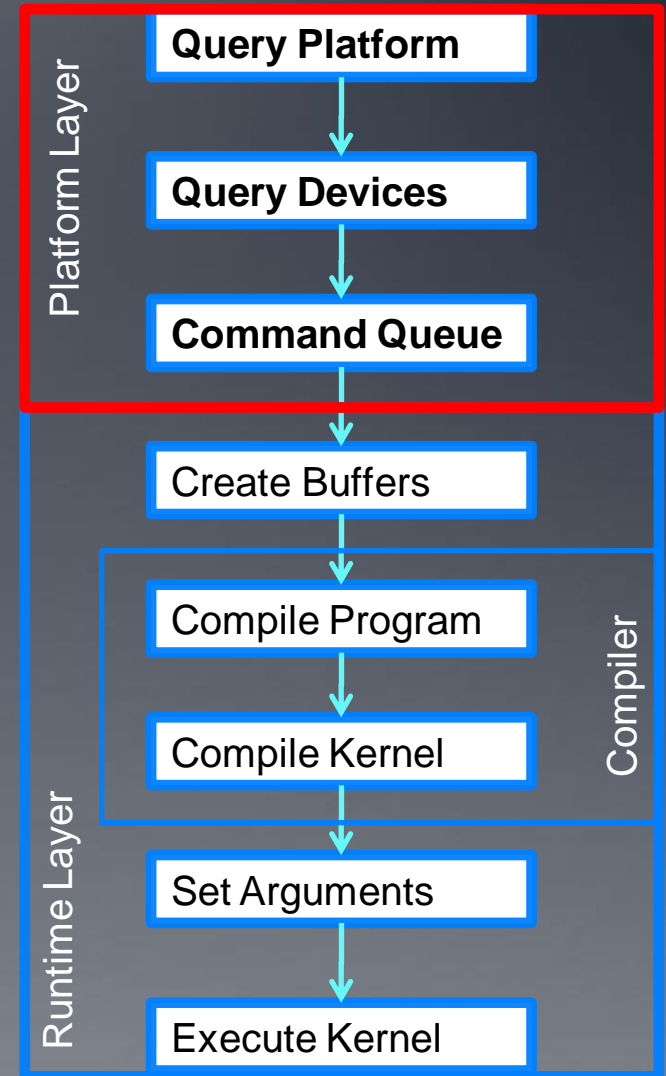
Step0: Initialize Device

- Declare context
- Choose a device from context
- Using device and context create a command queue

```
cl_context myctx = clCreateContextFromType (  
    0, CL_DEVICE_TYPE_GPU,  
    NULL, NULL, &ciErrNum);
```

```
ciErrNum = clGetDeviceIDs (0,  
    CL_DEVICE_TYPE_GPU,  
    1, &device, cl_uint *num_devices)
```

```
cl_commandqueue myqueue ;  
myqueue = clCreateCommandQueue(  
    myctx, device, 0, &ciErrNum);
```



Step1: Create Buffers

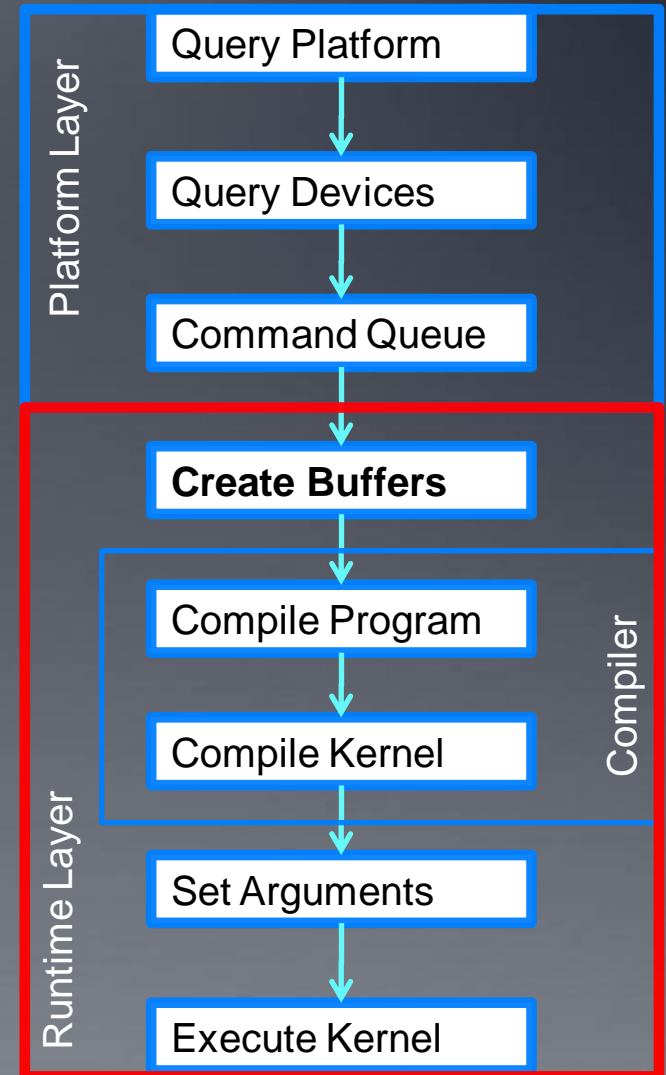
- Create buffers on device
 - Input data is read-only
 - Output data is write-only

```
cl_mem d_ip = clCreateBuffer(  
    myctx,  
    CL_MEM_READ_ONLY,  
    mem_size,  
    NULL, &ciErrNum);
```

```
cl_mem d_op = clCreateBuffer(  
    myctx,  
    CL_MEM_WRITE_ONLY,  
    mem_size,  
    NULL, &ciErrNum);
```

- Transfer input data to the device

```
ciErrNum = clEnqueueWriteBuffer (  
    myqueue , d_ip, CL_TRUE,  
    0, mem_size, (void *)src_image,  
    0, NULL, NULL)
```



Step2: Build Program, Select Kernel

// create the program

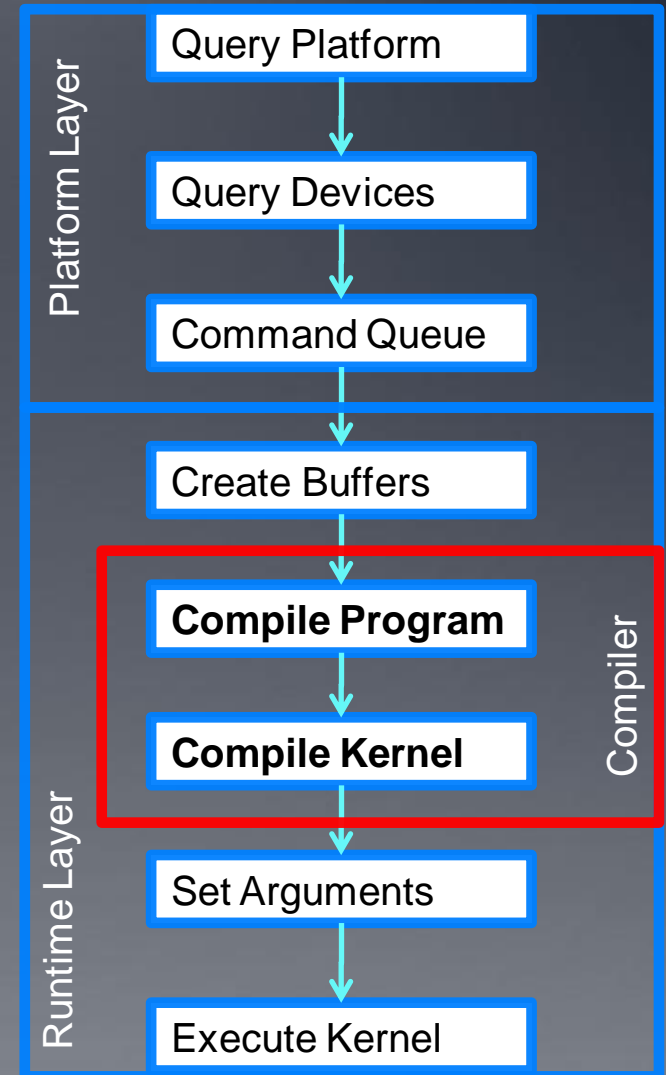
```
cl_program myprog = clCreateProgramWithSource  
( myctx,1, (const char **)&source,  
  &program_length, &ciErrNum);
```

// build the program

```
ciErrNum = clBuildProgram( myprog, 0,  
  NULL, NULL, NULL, NULL);
```

//Use the "image_rotate" function as the kernel

```
cl_kernel mykernel = clCreateKernel (  
  myprog , "image_rotate" ,  
  error_code)
```



Step3: Set Arguments, Enqueue Kernel

// Set Arguments

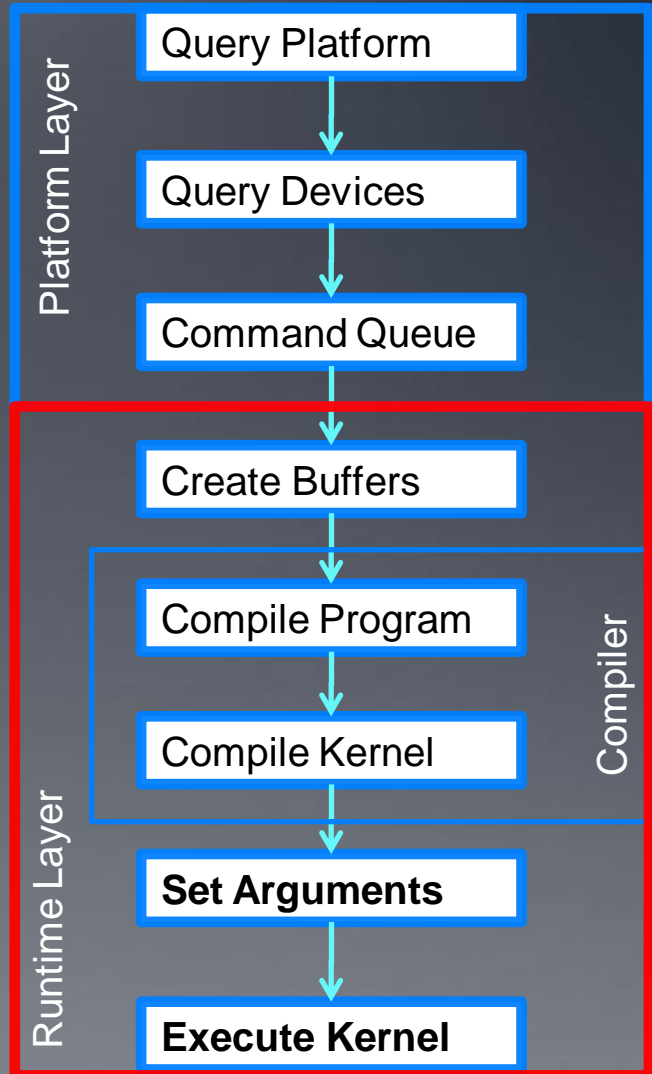
```
clSetKernelArg(mykernel, 0, sizeof(cl_mem),  
               (void *)&d_ip);  
clSetKernelArg(mykernel, 1, sizeof(cl_mem),  
               (void *)&d_op);  
clSetKernelArg(mykernel, 2, sizeof(cl_int),  
               (void *)&W);
```

//Set local and global workgroup sizes

```
size_t localws[2] = {16,16} ;  
size_t globalws[2] = {W, H}; //Assume divisible by 16
```

// execute kernel

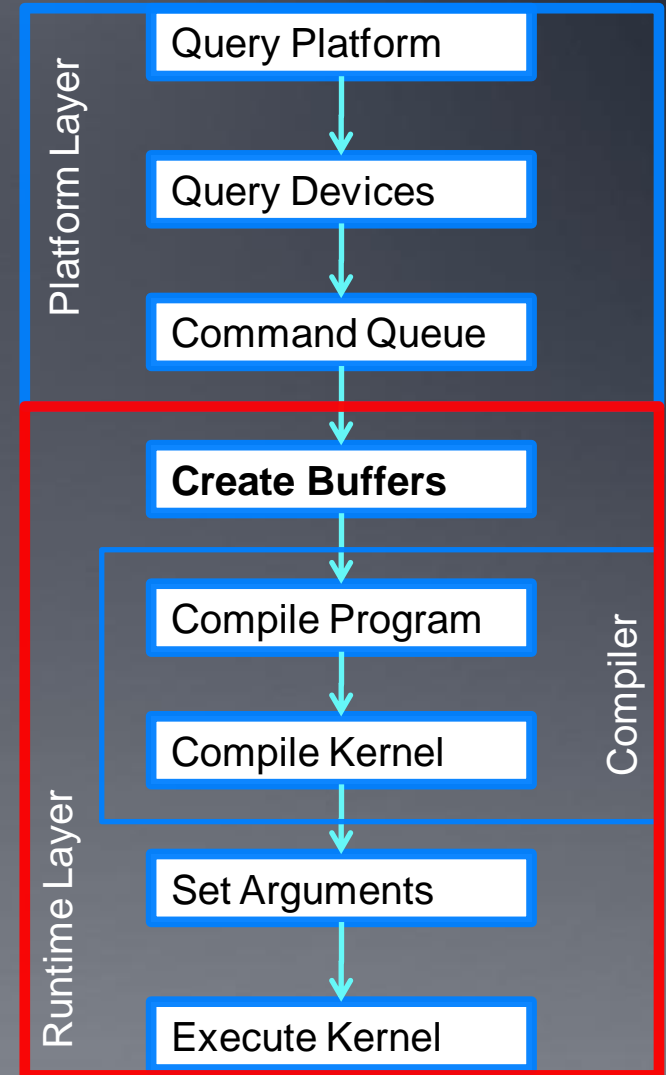
```
clEnqueueNDRangeKernel(  
    myqueue , myKernel,  
    2, 0, globalws, localws,  
    0, NULL, NULL);
```



Step4: Read Back Result

- Only necessary for data required on the host
- Data output from one kernel can be reused for another kernel
 - Avoid redundant host-device IO

```
// copy results from device back to host
clEnqueueReadBuffer(
    myctx, d_op,
    CL_TRUE,      //Blocking Read Back
    0, mem_size, (void *) op_data,
    NULL, NULL, NULL);
```



OpenCL Timing

- OpenCL provides “events” which can be used for timing kernels
 - Events will be discussed in detail in Lecture 11
- We pass an event to the OpenCL enqueue kernel function to capture timestamps
- Code snippet provided can be used to time a kernel
 - Add profiling enable flag to create command queue
 - By taking differences of the start and end timestamps we discount overheads like time spent in the command queue

```
cl_event event_timer;  
clEnqueueNDRangeKernel(  
    myqueue, myKernel,  
    2, 0, globalws, localws,  
    0, NULL, &event_timer);
```

```
unsigned long starttime, endtime;
```

```
clGetEventProfilingInfo(event_time,  
    CL_PROFILING_COMMAND_START,  
    sizeof(cl_ulong), &starttime, NULL);
```

```
clGetEventProfilingInfo(event_time,  
    CL_PROFILING_COMMAND_END,  
    sizeof(cl_ulong), &endtime, NULL);
```

```
unsigned long elapsed =  
(unsigned long)(endtime - starttime);
```

Example 2

Matrix Multiplication

Basic Matrix Multiplication

- Non-blocking matrix multiplication
 - Doesn't use local memory
 - Each element of matrix reads its own data independently
- Serial matrix multiplication

```
for(int i = 0; i < Ha; i++)  
    for(int j = 0; j < Wb; j++){  
        c[i][j] = 0;  
        for(int k = 0; k < Wa; k++)  
            c[i][j] += a[i][k] + b[k][j]  
    }
```

- Reuse code from image rotation
 - Create context, command queues and compile program
 - Only need one more input memory object for 2nd matrix

Simple Matrix Multiplication

```
__kernel void simpleMultiply(  
    __global float* c, int Wa, int Wb,  
    __global float* a, __global float* b) {
```

//Get global position in Y direction

```
int row = get_global_id(1);
```

//Get global position in X direction

```
int col = get_global_id(0);
```

```
float sum = 0.0f;
```

//Calculate result of one element

```
for (int i = 0; i < Wa; i++) {
```


```
    sum +=
```

```
        a[row*Wa+i] * b[i*Wb+col];
```

```
}
```

```
c[row*Wb+col] = sum;
```

```
}
```



Summary

- We have studied the use of OpenCL buffer objects
- A complete program in OpenCL has been written
- We have understood how an OpenCL work-item can be used to work on a single output element (seen with rotation and matrix multiplication)
 - While the previously discussed examples are correct data parallel programs their performance can be drastically improved
- Next Lecture
 - Study the GPU memory subsystem to understand how data must be managed to obtain performance for data parallel programs
 - Understand possible optimizations for programs running on data parallel hardware like GPUs